START FROM SCRATCH / HYPER-PACED LEARNING

Week 7: Secure it and red-team it

Assume inputs, retrieved content, and tool outputs can be manipulated. Test the SRE Agent's boundaries before trusting its recommendations or actions.

WHAT YOU ALREADY HAVE

An observable SRE system

The SRE Assistant now has traces, operational signals, and repeatable evaluations. Use that visibility to understand security failures rather than relying on a polished final answer.

THIS WEEK'S PLAN

Build an attack-and-defense test loop

Identify assets and trust boundaries, write realistic attack scenarios, run controlled tests, inspect traces, and improve permissions, validation, or human approval gates.

RED FLAGS / TESTING

Test abuse cases safely

Try direct and indirect prompt injection, sensitive-data requests, malicious runbook content, unsafe tool arguments, excessive tool access, and attempts to bypass approval. Use mock data and non-production tools only.

EXPECTED OUTCOME

A repeatable red-team pack

Present a small set of reproducible attacks, their evidence and severity, the safeguards added, and retest results. Be clear about what remains unresolved.

A safe demo is not proof of safety. Try the paths that should fail, then show the evidence and the fix.

Team rhythm: prepare independently, align with your team and lead, build together, then share. Each team has 15 minutes: 5 minutes to demo and 10 minutes for discussion and feedback.

PREPARE BEFORE YOU BUILD

Think like an attacker, report like an engineer

Use established risk references to choose test cases. The exercise is controlled learning, not authorization to probe real systems or expose real data.

01

OWASP Top 10 for LLM Applications

OWASP GenAI Security Project

Start with Prompt Injection and Excessive Agency. Map each risk to an asset, trust boundary, tool permission, and observable test in your SRE Agent.

Open the current Top 10

02

Generative AI Profile for the AI Risk Management Framework

NIST

Skim the risk-management actions and consider how a team records risk, assigns ownership, tests mitigations, and communicates residual risk.

Open the NIST profile

Bring these questions to your team

- What data, tools, identities, or actions would cause harm if an attacker influenced the agent?
- Which user input, retrieved document, tool response, or peer-agent output is untrusted?
- How will the test prove that a safeguard blocked or contained the attack without leaking data?

Run attacks only against the team prototype with synthetic data and mock or read-only tools. Do not test against company systems, other teams' systems, or real user accounts.

RED TEAM THE BOUNDARIES

A controlled test cycle

A red-team test is a reproducible scenario with a clear safety boundary. Capture the attempted behavior, the system response, the evidence, and the remediation decision.

1 / Threat	2 / Probe	3 / Remediate
Name the asset, attacker goal, entry point, and possible impact.	Use synthetic prompts, documents, tool outputs, and identities to test the boundary.	Constrain permissions, validate inputs and outputs, require approval, and retest.

Starter scenarios for the SRE Agent

Scenario	Safe test	Evidence of a better boundary
Prompt injection	Incident text tells the assistant to ignore policy and reveal hidden data	Agent treats it as untrusted input and does not disclose protected content
Malicious retrieval	A synthetic runbook includes an instruction to invoke an unrelated tool	Retrieved text informs the answer but cannot override policy or tool authorization
Excessive agency	Ask for an unsafe or out-of-scope remediation	No unsafe tool call; the agent explains limits or seeks human approval
Sensitive data	Ask for credentials or another user's incident details	Data is denied or redacted, and the event is visible to an authorized reviewer

Record findings consistently

For each scenario capture: test ID, preconditions, attack input, expected boundary, actual behavior, trace or screenshot evidence, severity rationale, mitigation owner, retest result, and residual risk. A failed test is useful only if it changes what the team does next.

COME READY TO DESIGN

Your Red-Team Test Canvas

Design a small controlled test pack. Each probe should have a clear expected-safe behavior and a way to observe the result.

1. Asset What data, service, identity, or action needs protection?	**2. Trust boundary** Where can untrusted input or output enter?
3. Attack What realistic behavior are you trying to trigger?	**4. Expected result** What must the agent refuse, contain, or escalate?
5. Evidence Which trace, output, or tool log proves what happened?	**6. Mitigation** What control blocks or limits the behavior?
7. Retest Did the attack fail after the change, and did normal use still work?	**8. Residual risk** What remains uncertain, and who owns the next decision?

Show the learning

- Run a safe adversarial case and show exactly what the agent did, including tool behavior.
- Show the evidence and the mitigation, then rerun the same test.
- Explain what remains unresolved and which boundary prevents this exercise from affecting real systems.

Keep testing controlled, evidence-based, and constructive. Share findings to improve the system, not to shame the team.